

Graphical Models. Assignment

Dmitry Adamskiy

November 18, 2021

1 Inference and Learning

1. **Two earthquakes** Solve exercise 1.22 from the textbook [5 marks for each subquestion, 20 marks in total].
 2. **Meeting scheduling** You are organising a trip to Scotland for your N friends. You booked the tickets on the Caledonian Sleeper train, departing from Euston station. Since you have all the tickets you decide that you need to meet at Euston station meeting point before the ticket gates. You all need to be there at 21:05 in order not to miss the train. However, it makes sense to ask people to come a bit earlier in case some of them are delayed. But how much earlier?
 - (a) Here is the model for the delays that you are going to use. You are always on time. Let D_i be a delay of your i -th friend. You assume that all D_i -s are independent and identically distributed. $P(D_i \leq 0) = 0.7$ (that is, with probability 0.7 your friend will come on time or earlier), $P(0 < D_i < 5 \text{ mins}) = 0.1$, $P(5 \text{ mins} \leq D_i < 10 \text{ mins}) = 0.1$, $P(10 \text{ mins} \leq D_i < 15 \text{ mins}) = 0.07$, $P(15 \text{ mins} \leq D_i < 20 \text{ mins}) = 0.02$, $P(20 \text{ mins} \leq D_i) = 0.01$. You would like to meet as late as possible but still catch a train with probability guaranteed to be at least 0.9. What time $T_0 = T_0(N)$ should you ask your friends to meet? Solve for $N = 3$, $N = 5$, $N = 10$. [5 marks]
 - (b) You realise that some people are less punctual than others. You update your model with the unobserved binary variables Z_i . The probabilities $P(D_i | Z_i = \text{punctual})$ are the same as above and $P(D_i | Z_i = \text{not punctual}) = (0.5, 0.2, 0.1, 0.1, 0.05, 0.05)$ where the states are ordered as above. You have a prior belief that $p(Z_i = \text{punctual}) = 2/3$ independently for all i . What are the probabilities of missing the train if you use the answers from (a) for this model? [5 marks]
- Bonus Suppose that for $N = 5$ you used the answer from (a) and missed the train. Now you are wondering how many of your N friends are not punctual. What is the posterior distribution of this count? [5 marks]

3. **Three weather stations.** You have the data from three weather stations (`meteo1.csv`) and you don't know which data sequences come from which one. You believe that the data for each weather station follow a first-order Markov chain.
- (a) Draw the graphical model for this problem. What are the parameters of this model? Describe all the steps of the EM-algorithm in this scenario. [5 marks]
 - (b) Implement the EM algorithm for the data provided. Print the learned parameters and the log-likelihood of the data for those parameters. For the first 10 rows of the dataset print the posterior distribution of which stations these data sequences come from. [5 marks]
 - (c) Explain whether the learned parameters depend on the initial guess for the parameters of the model. Explain your initialisation strategy. Describe any computational issues you encountered when implementing EM and how you solved them. [5 marks]
 - (d) Suppose that you don't have any prior knowledge about the parameters. Your friend tells you that you should then initialise all the relevant distributions to be uniform. Is this a good idea? What would the parameters learned by EM look like and why? [5 marks]

2 LDPC codes

2.1 Background: LDPC-codes

Here we briefly review the concepts and the construction of LDPC decoder and we invite you to study [1] and [2] for more details and motivations. All the addition and multiplication operations in what follows are operations in \mathbb{F}_2 (that is, $0+1=1+0=1, 1+1=0+0=0, 1\cdot 0=0\cdot 1=0\cdot 0=0$ and $1\cdot 1=1$).

The message to be transmitted is split into blocks of length K . Let's denote one such block to be transmitted as $t \in \{0, 1\}^K$. Let N be the codeword length. The code is a subset of 2^K codewords of length N and the ratio K/N is called a rate of the code. A *linear* code is a code such that all the codewords form a K -dimensional linear subspace of $\{0, 1\}^N$. Such subspace could be defined by a parity check matrix H or rank $M = N - K$. The codewords are then defined as the solutions of a system of linear equations $Hx = 0$. In what follows we assume that matrix H is full-rank of size $(N - K, N)$.

2.1.1 Encoding

Suppose we are given some parity check matrix H and we need to find a generator matrix G . If we know the basis vector of codeword subspace, we could use the basis vectors as the columns of matrix G . This way $x = Gt$ is going to be a codeword. The encoding is called systematic if all the bits of t are copied to the specific location of the transmitted message x (for example in the first K bits). Then reconstructing the signal from the decoded message becomes trivial: you just read it from the first K bits. One way to build a systematic encoder G is to perform Gaussian elimination. As described in the tutorial, up to the permutation of columns, the echelon form of H is equivalent to $[PI_{N-K}]$. Then you could select G to be $[I_K P]^T$, as $HGt = (I + I)t = 0$ for every t .

2.1.2 Decoding

The probabilistic model of LDPC decoder is as follows. Let x be a transmitted vector and y the received one. The noise model specifies conditional probability distributions $P(y|x)$. For example, in the Binary Symmetric Channel model each bit is independently flipped with probability p , so we have

$$P(y|x) = \prod_{n=1}^N p(y_n|x_n) = \prod_{n=1}^N p^{x_n-y_n}(1-p)^{x_n-y_n+1}$$

The joint distribution of (x, y) is then defined as $p(x, y) = p(y|x)p(x)$, where $p(x)$ is a uniform prior distribution over all the valid codewords:

$$p(x) \propto \mathbb{I}[Hx = 0]$$

Decoding is done using Loopy Belief Propagation as described in the tutorial slides and in [1, p.p. 560-561].

2.2 Assignment

1. Write a function that receives a parity check matrix H and builds a systematic encoding matrix G for it. This may require renaming the variables, so

the function should return two matrices: \hat{H} and G , such that \hat{H} is equal to H up to a column permutation and $\hat{H}Gt = 0$ for all t (all the operations are performed in \mathbb{F}_2). Print the outputs of the function for the following matrix (note: we may test it on other inputs as well!):

$$H = \begin{pmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 \end{pmatrix}$$

[8 marks]

2. Draw the factor-graph for the same matrix. Write the distribution corresponding to this factor-graph and the updates used for the messages. [2 marks]
3. Write an LDPC-decoder based on Loopy Belief Propagation for Binary Symmetric Channel. Specifically, write a function that receives a parity check matrix \hat{H} , a received word y , a noise ratio p and an optional parameter of a maximum number of iterations (with default value of 20). The function should return a decoded vector along with the following return code: 0 for success, -1 if the maximum number of iterations is reached without a successful decoding. Try to make your code efficient. Print the result of the decoding for a given parity check matrix H_1 (in H1.txt) and vector y_1 (in y1.txt). The noise ratio was $p = 0.1$. How many iterations did your algorithm take to converge? [12 marks]
4. The original message is located in the first 252 bits of the decoded signal. Recover the original English message by reading off the first 248 bits of the 252-bit message and treating them as a sequence of 31 ASCII symbols. [3 marks]

2.3 Note

1. You are not allowed to include in your program any code from existing LDPC-related packages.

3 Mean Field Approximation and Gibbs Sampling

Consider the Ising model on the $n \times n$ lattice as in Exercise 6.7 from [3] with the potentials modified to include a temperature-like parameter β : $P(x) = Z^{-1} \prod_{i>j} \phi(x_i, x_j)$ with $\phi(x_i, x_j) = e^{\beta \mathbb{I}[x_i=x_j]}$ for i a neighbour of j on a lattice and $i > j$ (to avoid overcounting).

3.1 Assignment

You will need to compute the joint probability distribution of the top and bottom nodes of the rightmost column of the 10×10 lattice. If $x_{i,j}$ is the node in i -th row and j -th column, that would be nodes $x_{1,10}$ and $x_{10,10}$, so you need to provide the probability table for $P(x_{1,10}, x_{10,10})$. You have to do it for the three values of β : $\beta = 4$, $\beta = 1$ and $\beta = 0.01$. For each of them, you have to do

it in the following three ways, printing the resulting probability distribution for each of them.

1. Perform exact inference, using techniques from Exercise 6.7. That is, treat each column as one variable with 2^n states and perform message passing on the induced factor-graph. [5 marks]
2. Use Mean Field Approximation and coordinate ascent. [10 marks]
3. Use Gibbs Sampling. [10 marks]

For each of the methods, write in your report the description of the methods and all the update equations used.

References

- [1] David J. C. MacKay. *Information Theory, Inference & Learning Algorithms*. Cambridge University Press, New York, NY, USA, 2002.
- [2] Amin Shokrollahi. LDPC codes: An introduction. 2002.
- [3] D. Barber. *Bayesian Reasoning and Machine Learning*. Cambridge University Press, 2012.